

Canada Median Income

Minh Tran, Fahad Bhutta, Aktan Kenzhebay, Saba Abbasi

2025-11-19

Abstract

This statistical study aims to determine patterns and trends in the median income of people living in Canada at two time periods: 2015 and 2020. The data used in the study was retrieved from Statistics Canada.

Through this study, we determined the difference (or lack thereof) in household income in specific scenarios such like time periods and household sizes/situations. From these findings, we hope to be able to make predictions and recommendations to relevant parties.

Introduction

The dataset chosen is Household Income Statistics by household type, and it includes households from all provinces and territories in Canada. The data was collected through a census, and it has data from broad, general areas like whole provinces, and smaller, more specific areas like cities. The census in which data was derived from was done in 2016 and 2021, and the median incomes provided was from 2015 and 2020 respectively.

Link to dataset: <https://www150.statcan.gc.ca/t1/tbl1/en/tv.action?pid=9810005701>

Link to cleaned dataset: https://drive.google.com/file/d/1kKOXK_0hBLu9F0nAYekCuXURS9noOP5v/view?usp=drive_link

Research questions

Question 1: Is there a difference in median after tax household income between 2015 and 2020?

Question 2: Is there a significant difference in income between households of 5 or more people and 4 persons (in 2020)? Do households with 5 or more people earn more than households with 4 people?

Question 3: Is there a difference in median income between Households with and without children? Do households with children earn more than households without children?

Variables selected:

Question 1: Median household after-tax income (2015), Median household after-tax income (2020)

Question 2: Median household total income (2020), Household Size: ‘4 persons’ and ‘5 or more persons’

Question 3: Median household total income (2020), Household type including census family structure: ‘With children’ and ‘Without children’

We excluded any rows that were had empty cells (missing values) or had a ‘0’ in the Number of Households column. Since there are times when the column had data for 2016 and not 2021, or vice versa, we chose to completely remove any rows that had empty cells.

Statistical Results

Question 1: Is there a difference in median after tax household income between 2015 and 2020?

Introduction

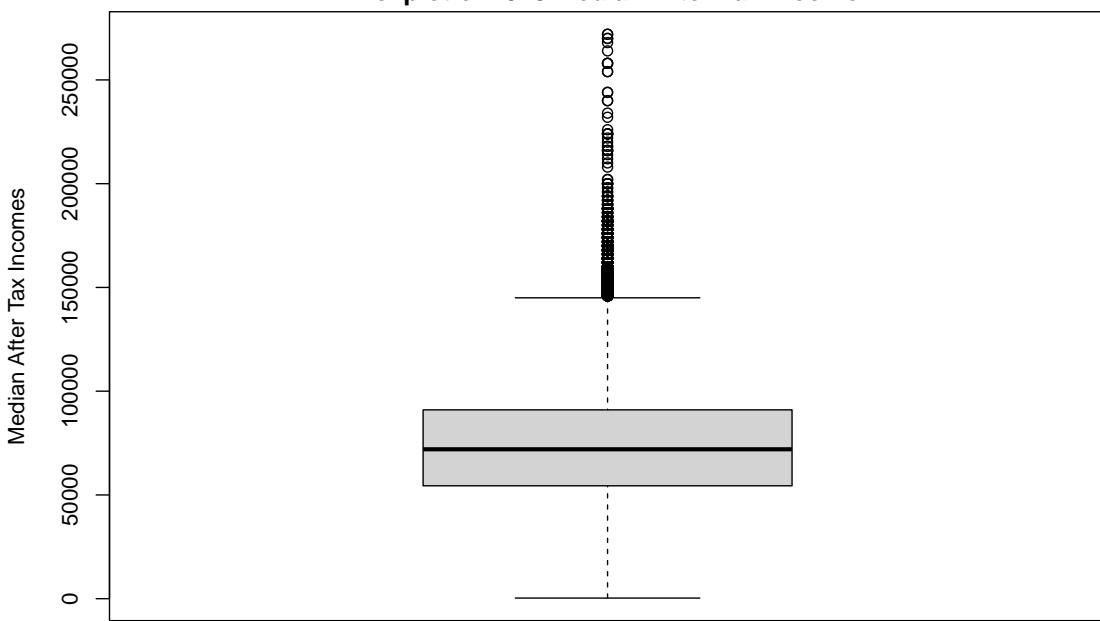
To explore how household economic conditions have changed over time, we examine whether median after tax household income differs between 2015 and 2020. This question helps identify broader trends in income growth and provides insight into how households may have been affected by economic shifts during this five-year period. Using the available income data for both years, we compare their distributions to determine if a meaningful difference exists.

min	Q1	median	Q3	max	mean	sd	n	missing
7850	62800	79500	1e+05	3e+05	81984.76	26934.17	156189	0

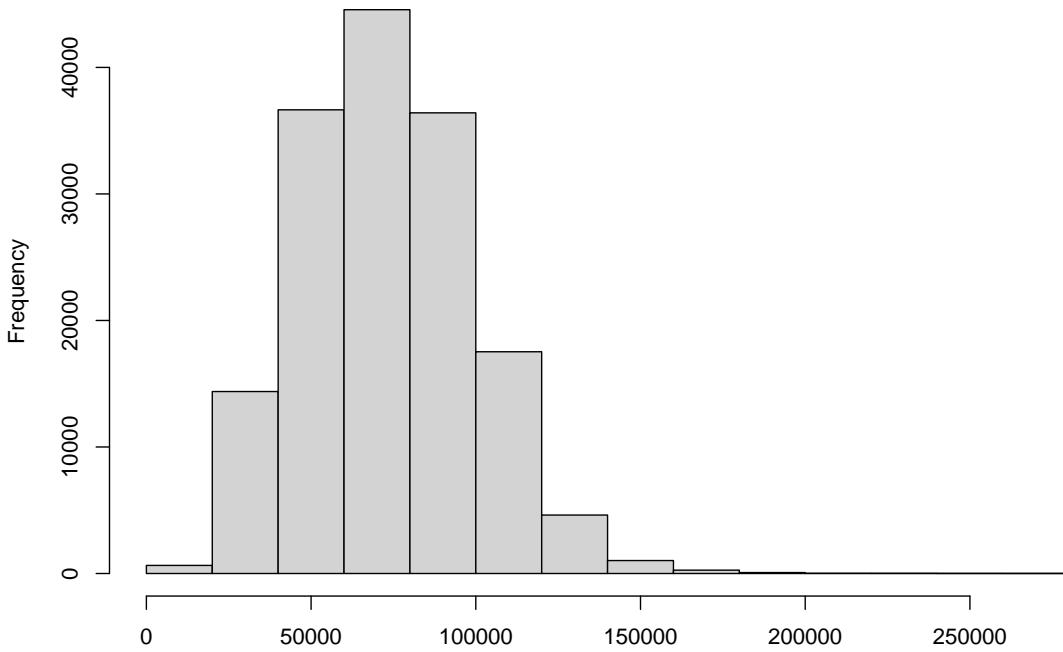
min	Q1	median	Q3	max	mean	sd	n	missing
294	54400	72000	91000	272000	73357.16	26053.9	156189	0

The summary statistics indicate that 2020 median after-tax household incomes are higher across all quartiles compared to 2015. The median increased from \$72,000 in 2015 to \$79,500 in 2020, and the mean increased from \$73,357 to \$81,985, suggesting an overall upward shift in income levels. These differences motivate a formal hypothesis test to determine whether the increase is statistically significant.

Boxplot of 2015 Median After Tax Income

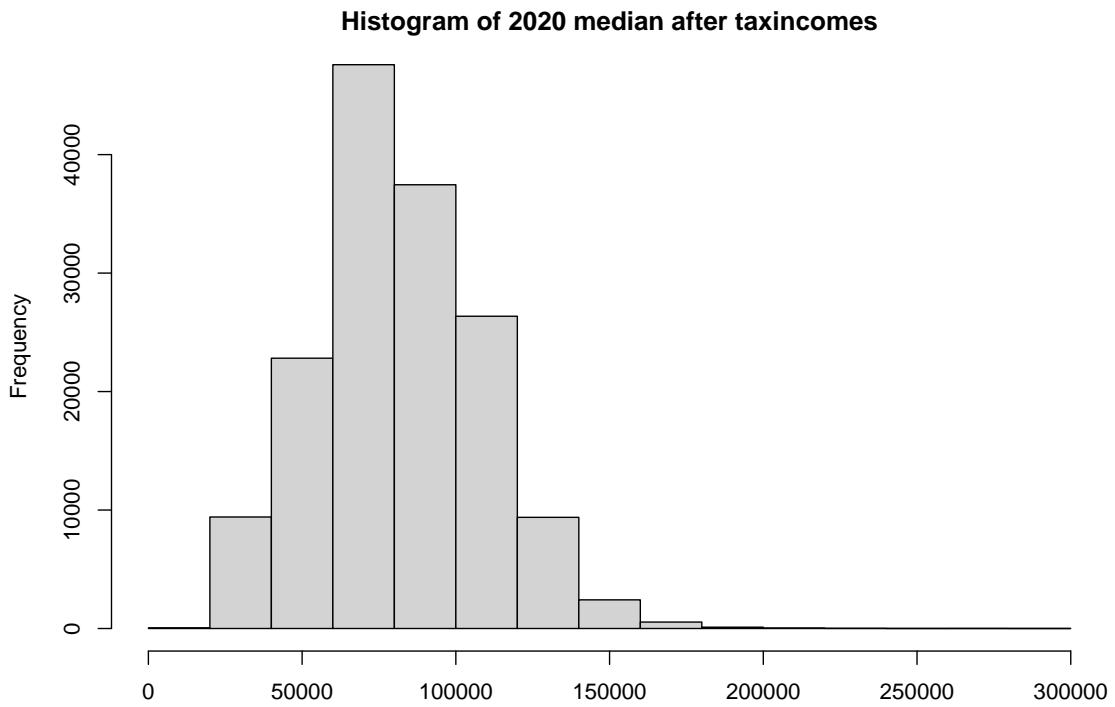
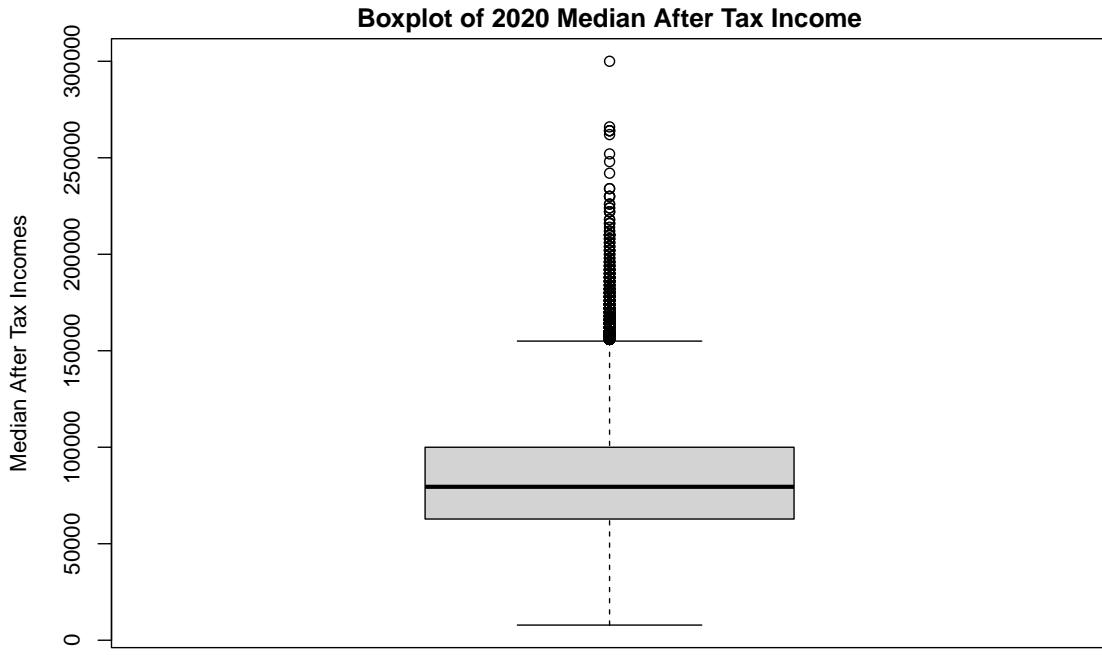


Histogram of 2015 median after tax incomes



The boxplot shows that the distribution of 2015 median after-tax incomes is right-skewed, with many high-income outliers extending far above the upper quartile. Most households fall within a moderate income range, with the middle 50% of values concentrated around roughly \$50,000–\$90,000.

The histogram similarly confirms a strong right skew: the majority of households earn between \$40,000 and \$100,000, while relatively few households have extremely high incomes. Together, these plots indicate that in 2015, median household incomes were unevenly distributed, with a long tail of higher-income households pulling the distribution to the right.



The 2020 boxplot shows a right-skewed distribution of median after-tax incomes, with many high-income outliers extending far above the upper quartile. Most households fall within a higher income range compared to 2015, with the middle 50% concentrated around roughly \$60,000–\$100,000.

The histogram also displays a pronounced right skew: the majority of households earn between \$50,000 and \$110,000, while a smaller number earn substantially more. Overall, these plots indicate that in 2020, household incomes increased relative to 2015 but still remained unevenly distributed with a long upper tail.

We performed a Welch's two sample t-test since the population variance is unknown, with Null Hypothesis (H_0) being the median is equal between 2015 and 2020 ($\text{inc2020} = \text{inc2015}$), and the Alternative Hypothesis (H_1) being the medians are not equal ($\text{inc2020} \neq \text{inc2015}$).

The p-value outputted from the t-test is less than 2.2e-16, which is extremely small, much smaller than our significance level of 5%. Thus, we conclude that there is a difference in the Median After Tax Household Income between 2015 and 2020.

Additionally, the confidence interval generated for the difference between 2020 and 2015 median after tax income ($\text{inc2020} - \text{inc2015}$) is (8441.755, 8813.442), which has positive upper and lower confidence limits. This suggests that we can be 95% confident in concluding that the Median Household After Tax Income is higher in 2020 than 2015.

The statistical results show that 2020 median after-tax household incomes are significantly higher than those in 2015, with a confidence interval indicating an average increase of roughly **\$8,400-\$8,800**. If similar economic and demographic patterns persist, this upward shift suggests that median household incomes in Canada may continue to rise in future census periods. However, this extrapolation should be made cautiously, because the dataset reflects regional medians, not individual households, and therefore generalizes trends rather than capturing precise population-level income growth.

Question 2: Is there a significant difference in income between households of 5 or more people and 4 persons (in 2020)?

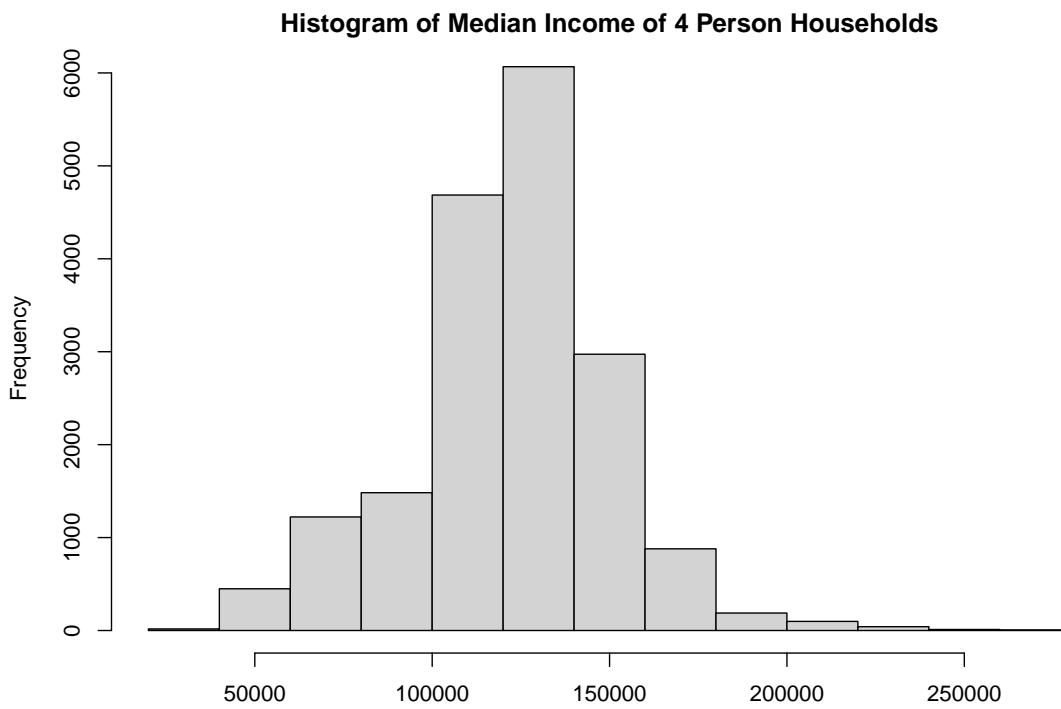
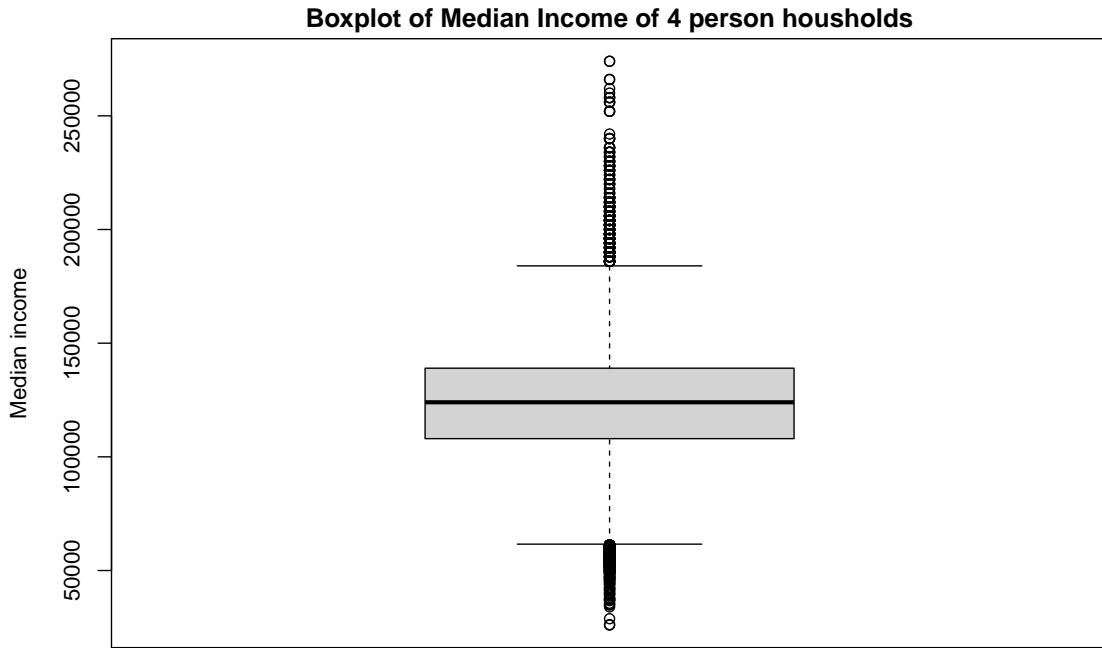
Introduction

To better understand how household size relates to income levels, we examine whether median after-tax household income differs between households with four people and those with five or more people in 2020. Comparing these two groups allows us to assess whether larger households tend to have higher or lower incomes and whether family size is associated with meaningful differences in financial resources. Using the available 2020 data, we compare the income distributions of both household sizes to determine if a significant difference exists.

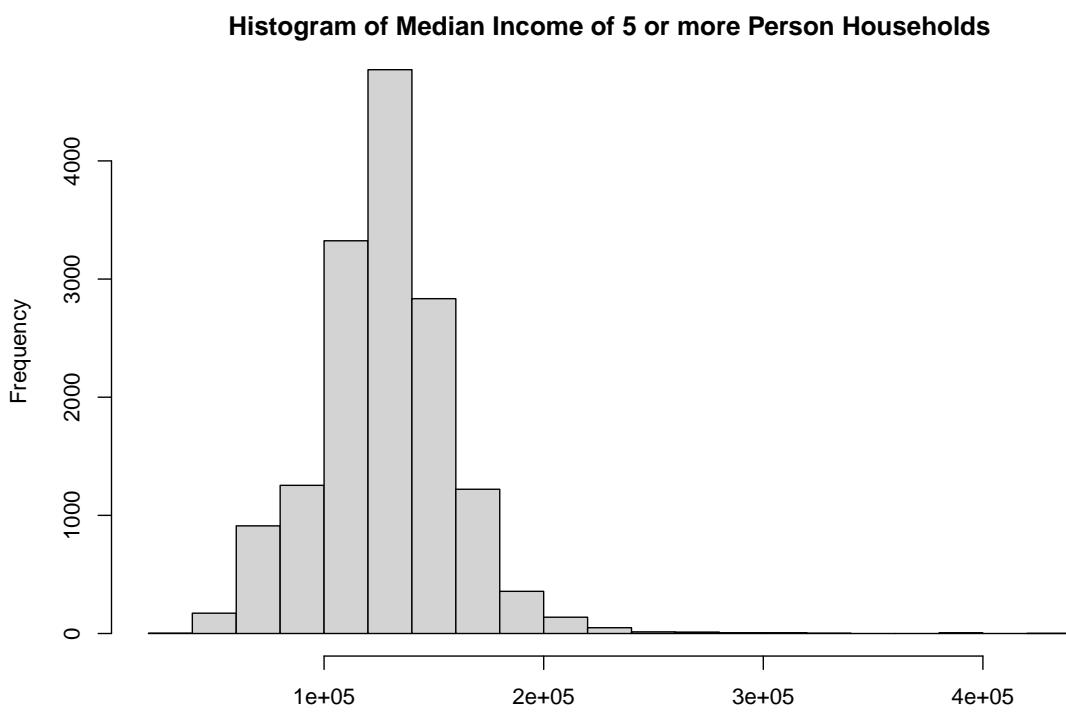
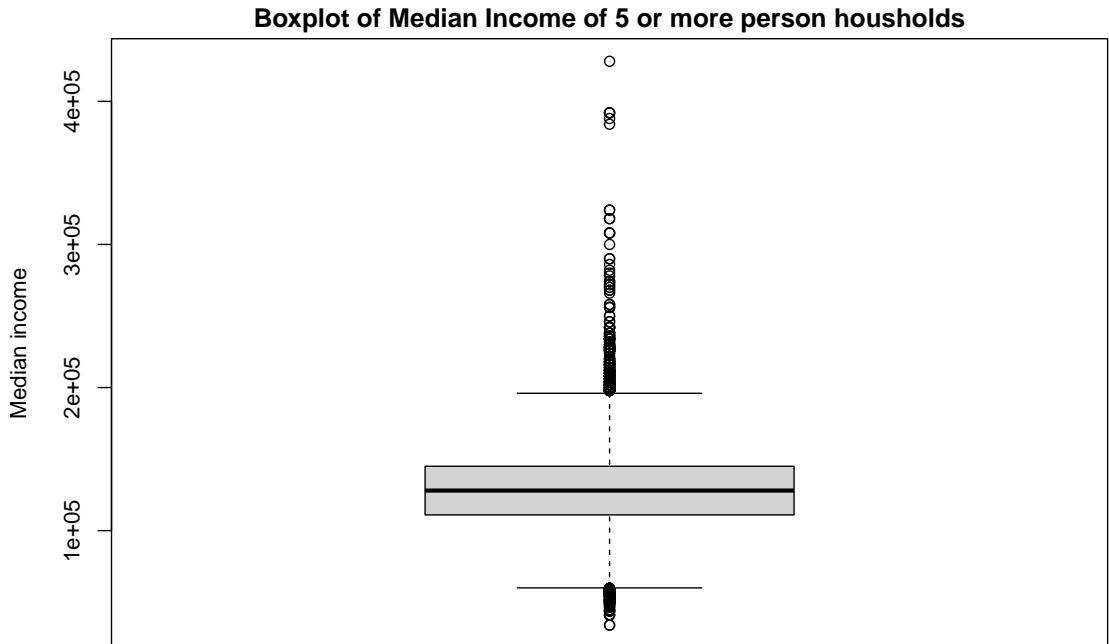
min	Q1	median	Q3	max	mean	sd	n	missing
26000	108000	124000	139000	274000	122603.8	28411.19	18119	0

min	Q1	median	Q3	max	mean	sd	n	missing
34000	111000	128000	145000	428000	128213.6	30803.87	15078	0

The summary statistics show that households with 5 or more people have consistently higher incomes than 4-person households. The median income increases from \$124,000 (4-person) to \$128,000 (5+), and the mean also rises from \$122,604 to \$128,214. Both groups show wide variability, but the shift upward across the quartiles suggests that larger households generally report higher incomes. These descriptive differences motivate a formal Welch t-test to determine if the income gap is statistically significant.



The boxplot for 4-person households shows a right-skewed income distribution, with several high-income outliers extending well above the upper quartile and most households concentrated around the \$100,000–\$150,000 range. The histogram supports this pattern, showing that the majority of 4-person households earn within this central band, while relatively few households have either very low or very high median incomes. Overall, the graphs indicate that although incomes for 4-person households tend to be relatively high, the distribution remains uneven, with a long upper tail reflecting a smaller group of substantially higher-earning households.



The boxplot for households with 5 or more people shows a right-skewed distribution of median incomes, with numerous high-income outliers extending well above the upper quartile. Most households are concentrated in a higher income band, roughly between \$120,000 and \$180,000, indicating that larger households tend to have relatively strong median earnings. The histogram confirms this pattern, with the majority of households earning around \$110,000–\$170,000 and far fewer households at very low or extremely high income levels. Overall, these graphs suggest that 5-plus-person households generally have higher median incomes, but the distribution remains uneven with a noticeable upper tail.

We performed a Welch's two sample t-test since the population variance is unknown. The Null Hypothesis (H_0) is The median income between households of 4 persons and 5 or more persons are the same ($\text{inc4} = \text{inc5p}$), and the Alternative Hypothesis (H_1) is The median income are different.

The p-value we got was less than 2.2e-16, which is extremely small, much smaller than our significance level of 5%. Thus, we reject H_0 , and conclude that the Median Income between Households with 4 persons and 5 or more persons is different.

Additionally, the 95% confidence interval for the difference between Median Income of 4 person and 5 plus person households ($\text{inc4} - \text{inc5p}$) is (-6,252.363, -4,967.192). This indicates that we can be 95% confident that households with 4 persons earn less than household with 5 or more persons.

Households with 5 or more people show significantly higher median income than 4-person households, with the difference estimated at **\$5,000-\$6,200**. This pattern implies that larger households may generally have greater earning capacity, possibly due to multiple earners or different economic structures within larger families. If this relationship holds across other regions and years, areas with a greater proportion of large households may also exhibit higher overall household income levels. Nevertheless, because the dataset uses aggregated regional medians, we should not assume this applies to every individual household.

Question 3: Is there a difference in median income between Households with and without children?

Introduction

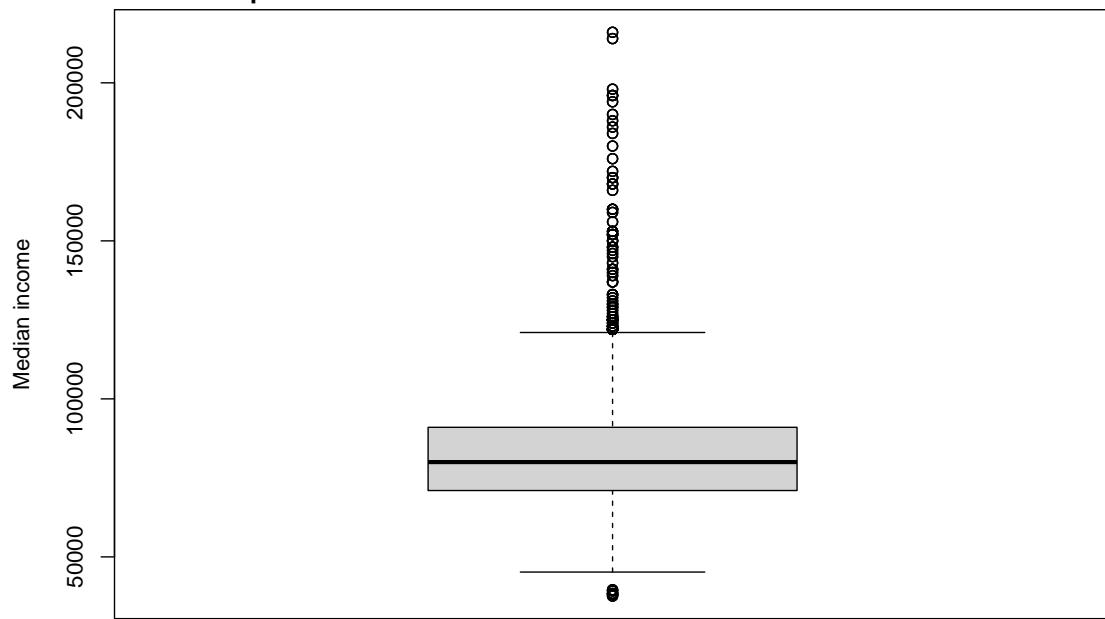
To understand how family structure relates to economic well being, we examine whether households with children differ in their median after tax income compared to households without children. This question helps reveal whether having children is associated with higher or lower household income levels and provides insight into how family composition may influence financial resources. Using the 2020 household income data, we compare the distributions for both groups to determine if a meaningful difference exists.

min	Q1	median	Q3	max	mean	sd	n	missing
37600	71000	80000	91000	216000	82239.62	17383.4	10965	0

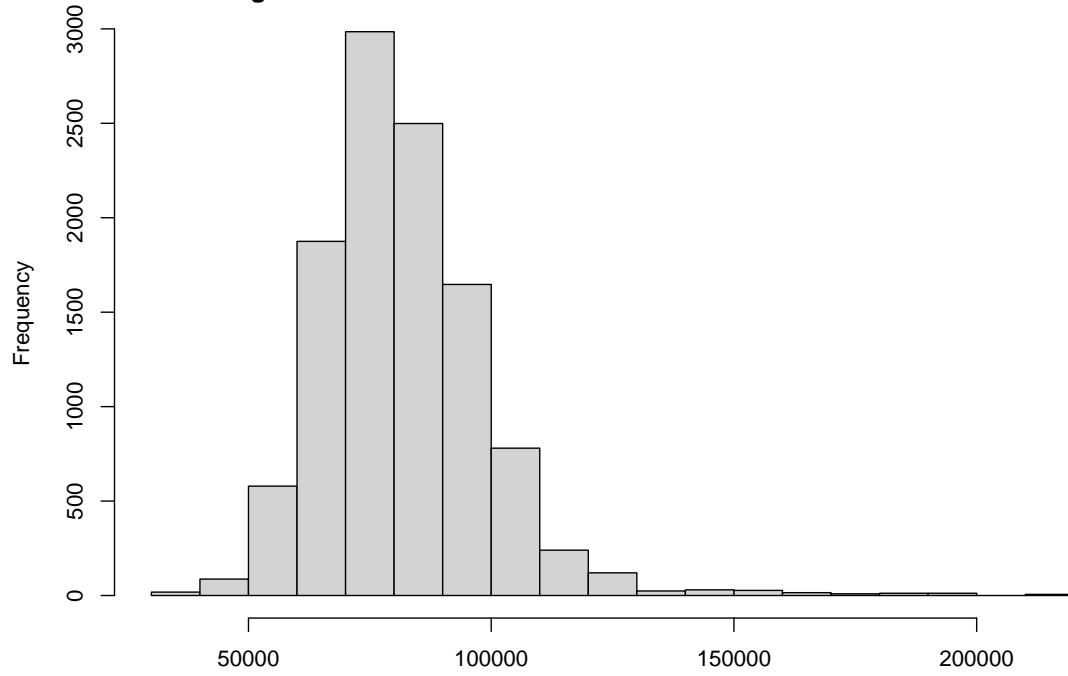
min	Q1	median	Q3	max	mean	sd	n	missing
36400	109000	121000	135000	392000	122754.9	23401.19	15685	0

The descriptive statistics indicate a large income difference between households with and without children. Households without children have a median income of \$80,000, while those with children show a much higher median of \$121,000. The means show a similar pattern, increasing from \$82,240 to \$122,755. Quartiles for households with children are also substantially higher, suggesting a strong upward shift in income. These differences lead us to use a Welch t-test to assess whether the income gap is statistically significant.

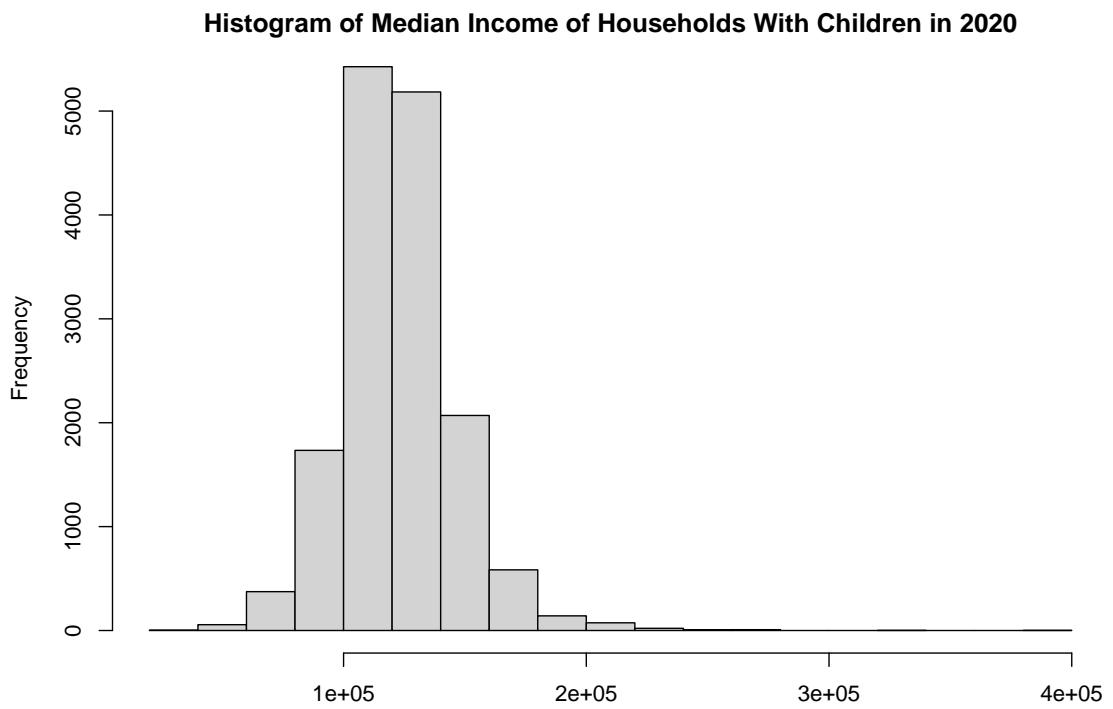
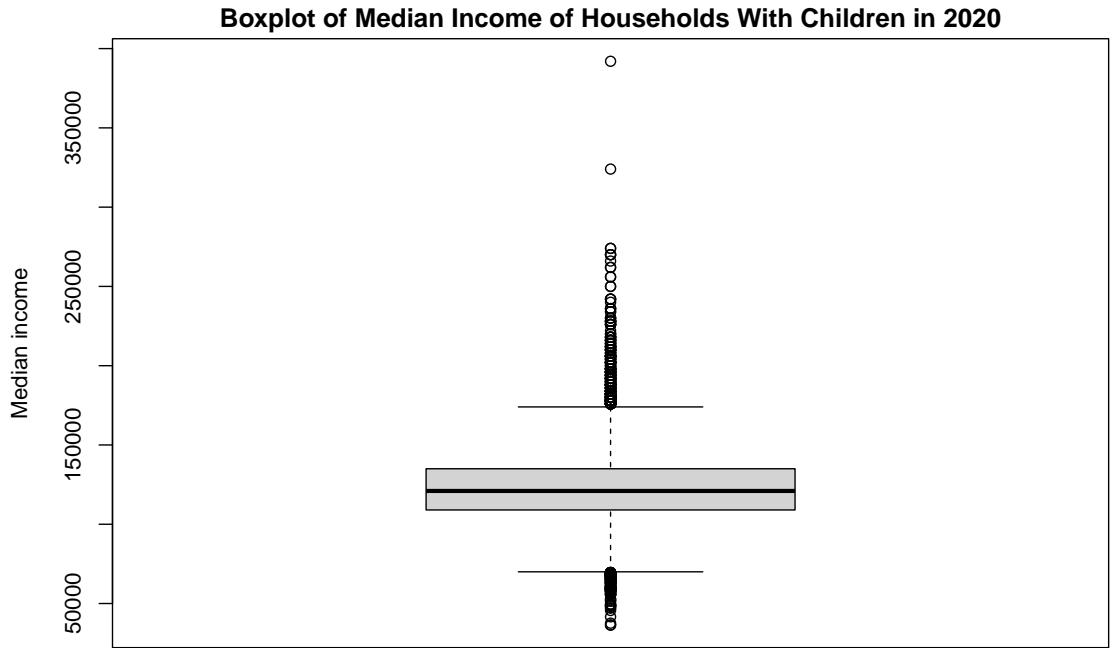
Boxplot of Median Income of Households Without Children in 2020



Histogram of Median Income of Households Without Children in 2020



The boxplot for households without children in 2020 shows a right-skewed distribution of median incomes, with several high-income outliers extending far above the upper quartile. Most households fall within a central range of roughly \$70,000–\$100,000, indicating relatively consistent earnings across this group. The histogram confirms this pattern, with the majority of households earning between about \$60,000 and \$110,000, while fewer households appear at very low or very high income levels. Overall, these plots suggest that although incomes for households without children cluster within a fairly narrow band, the distribution remains uneven with a noticeable upper tail of higher-earning households.



The boxplot for households with children in 2020 shows a right-skewed income distribution, with several high-income outliers far above the upper quartile. Most households cluster within a central range of roughly \$110,000–\$150,000, suggesting higher typical earnings compared to households without children. The histogram supports this pattern, with the majority of households earning between about \$100,000 and \$170,000, while far fewer appear at very low or very high income levels. Overall, these graphs indicate that households with children tend to earn higher median incomes, though the distribution still exhibits an uneven shape with a pronounced upper tail.

We performed a Welch's two sample t-test since the population variance is unknown. The Null Hypothesis (H_0) is there is no difference in median income between households with and without children ($\text{incW}_0 = \text{incW}$) and the Alternative Hypothesis (H_1) is the median household incomes are different ($\text{incW}_0 \neq \text{incW}$).

The p-value is less than 2.2-e16, which is smaller than our significance level at 5%. Thus, we can conclude that the Median Household Income between Families with children and without children is different.

The 95% confidence interval for the difference between the median income ($\text{incW}_0 - \text{incW}$) is (-41,005.23, -40025.42). Since both confidence limits are negative, we can conclude with 95% confidence that the median income of families who have children are higher than those without.

The results indicate that households with children earn substantially more than households without children, with the median income difference ranging from **\$40,025 to \$41,005**. This suggests that families who choose to have children tend to have higher income levels or greater financial capacity, which may reflect underlying factors such as number of earners, career stage, or long-term financial planning. If these conditions remain consistent, regions with more families may continue to show higher median incomes compared to regions dominated by child-free households. This inference should, however, be interpreted carefully because the dataset reports **regional medians**, not individual household income data.

Conclusion and Future Direction

Summary

1. 2015 vs 2020 after-tax household income

- The t-test shows a highly significant difference between 2015 and 2020 ($p < 2.2e-16$).
- The 95% CI for (2020 - 2015) is (8441.755, 8813.442), entirely positive.
- **Conclusion:** Median after-tax household income is higher in 2020 than in 2015.

2. Household size: 4 persons vs 5+ persons (2020)

- The t-test again shows a highly significant difference ($p < 2.2e-16$).
- The 95% CI for (4-person - 5+-person) income is (-6,252.363, -4,967.192), entirely negative.
- **Conclusion:** Households with 5 or more people earn more on average than 4-person households.

3. Households with vs without children

- The t-test finds a highly significant difference in income ($p < 2.2e-16$).
- The 95% CI for (without children - with children) is (-41,005.23, -40025.42), all negative.
- **Conclusion:** Households with children earn substantially more than households without children.

Limitations

1. One limitation is the dataset gives median incomes by region and household category, not individual households. Performing a t-test on median data would mean we are comparing the mean of medians, not the actual median. Hence, p-values and confidence intervals provide an approximation of the difference, not the true difference in population medians.
2. Another limitation is the dataset didn't have information for every single geographical location, which led to some rows either having 0 for Number of Households, and missing values for median incomes. One problem this could lead to is misrepresentation, since lower income regions could be missing, making my data more biased towards higher income areas
3. Another limitation is this dataset did not have information on key characteristics that affect earning like education, age, employment status, number of earners. These characteristics could potentially improve the precision of the analysis.

Possible extensions/improvements.

1. Get raw, individual data instead of grouped data. This will allow us to perform a much more accurate test, and the result will represent the actual mean better.
2. Unrealistic, but to get data from every region, and from every household. This will allow us to get accurate information on means and medians of each groups, which will provide an unskewed result that is representative of the whole population.
3. Get more characteristics from each sampled household, so we can better determine trends and correlations.

Appendix (R code)

Note that each question is done in a seperate R script, will be provided upon request.

Question 1: Is there a difference in median after tax household income between 2015 and 2020?

Load libraries

```
library(dplyr)
library(ggplot2)
```

Read data

```
income <- read.csv("incomes_cleaned.csv", check.names = FALSE)
```

filter by the first 11 rows of each geographic region that is not Canada

11 since these 11 rows have data of all household sizes

The difference between these rows are household types, e.g. Couple with children, without children, one parent, two parent, etc

```
income_sub <- income%>%
  filter(GEO != 'Canada') %>%
  filter(!is.na(`Household type including census family structure (11)`))
```

exclude any rows where income is 0

```
income_sub <- income%>%
  filter(
    !is.na(`Median household after-tax income (2020)`),
    !is.na(`Median household after-tax income (2015)`),
    `Median household after-tax income (2020)` != 0,
    `Median household after-tax income (2015)` != 0
  )
```

getting the data from their columns

```
inc2020 <- income_sub$`Median household after-tax income (2020)`
inc2015 <- income_sub$`Median household after-tax income (2015)`
```

Summary Statistics & variances

```
var_median_2020 <- var(inc2020)
var_median_2015 <- var(inc2015)

summary(inc2020)

##      Min. 1st Qu. Median      Mean 3rd Qu.      Max.
##    7850    62800   79500    81985 100000  300000

summary(inc2015)
```

```
##      Min. 1st Qu. Median      Mean 3rd Qu.      Max.
##    294     54400   72000    73357  91000  272000
```

```
var_median_2020
```

```
## [1] 725449693
```

```
var_median_2015
```

```
## [1] 678805957
```

Two sample t-test (unequal variances)

```
test <- t.test(inc2020, inc2015, alternative = "two.sided", var.equal = FALSE)
test
```

T-test since we do not know the population variance

```
##
##  Welch Two Sample t-test
##
## data: inc2020 and inc2015
## t = 90.99, df = 312032, p-value < 2.2e-16
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
##  8441.755 8813.442
## sample estimates:
## mean of x mean of y
## 81984.76 73357.16

p_val <- test$p.value
p_val
```

```
## [1] 0
```

Since P-value is smaller than 0.05, we reject H₀ at 5% significance, and conclude that there is a difference in median after tax household income between 2020 and 2015.

Assumptions

Data is continuous:

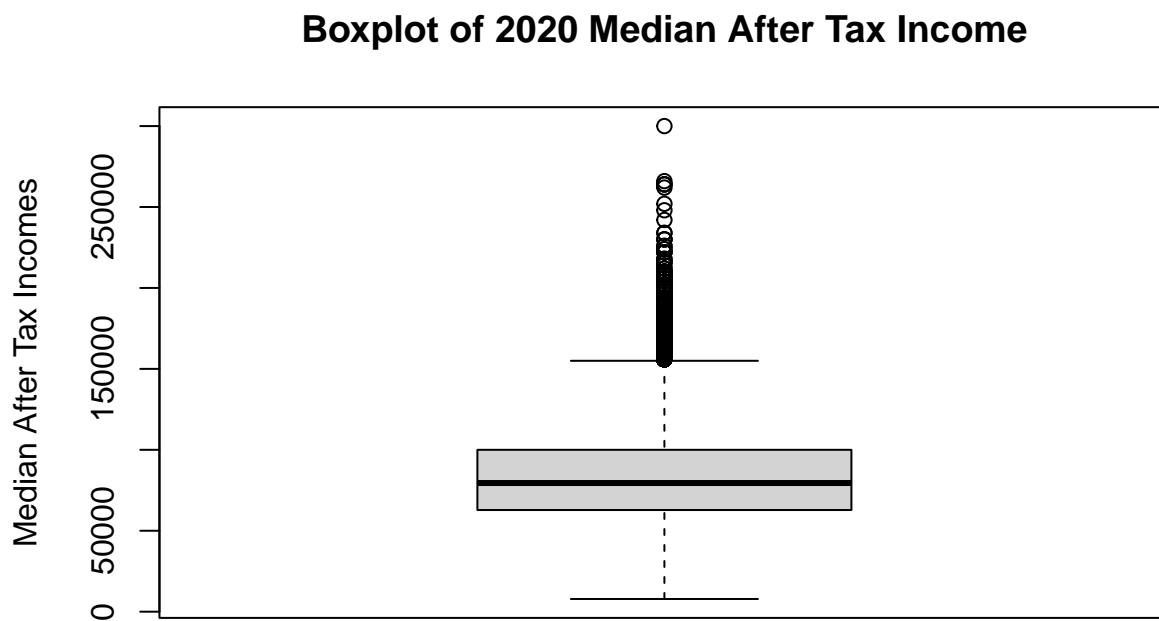
Data is independent:

This one is probably true since the data are from two different time periods

Distribution of data is approximately normal:

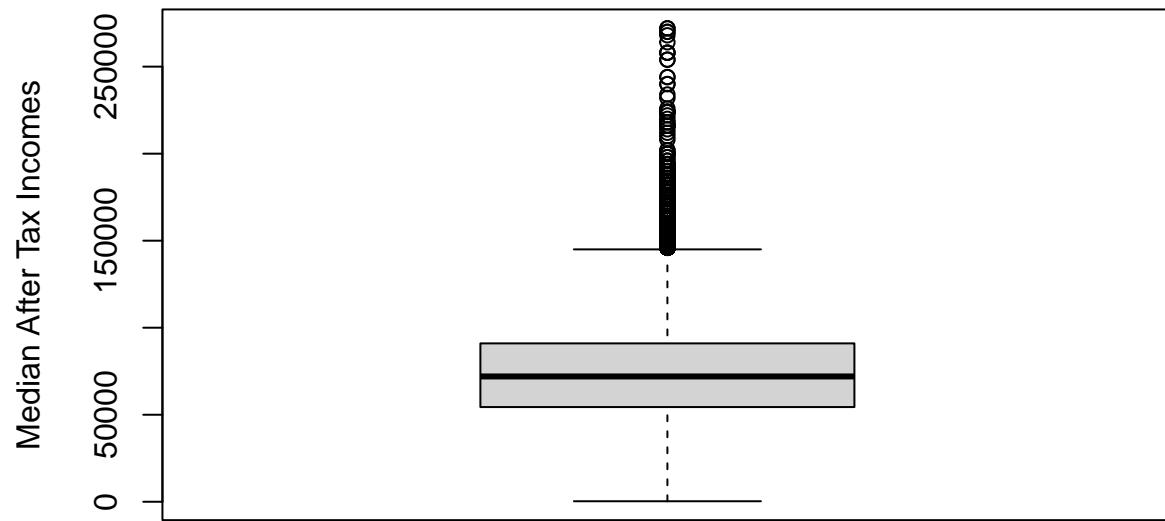
```
# Boxplots
box2020 = boxplot(inc2020, ylab = "Median After Tax Incomes", main = "Boxplot of 2020 Median After Tax Income")
```

We use boxplot, histogram and QQ plot for this



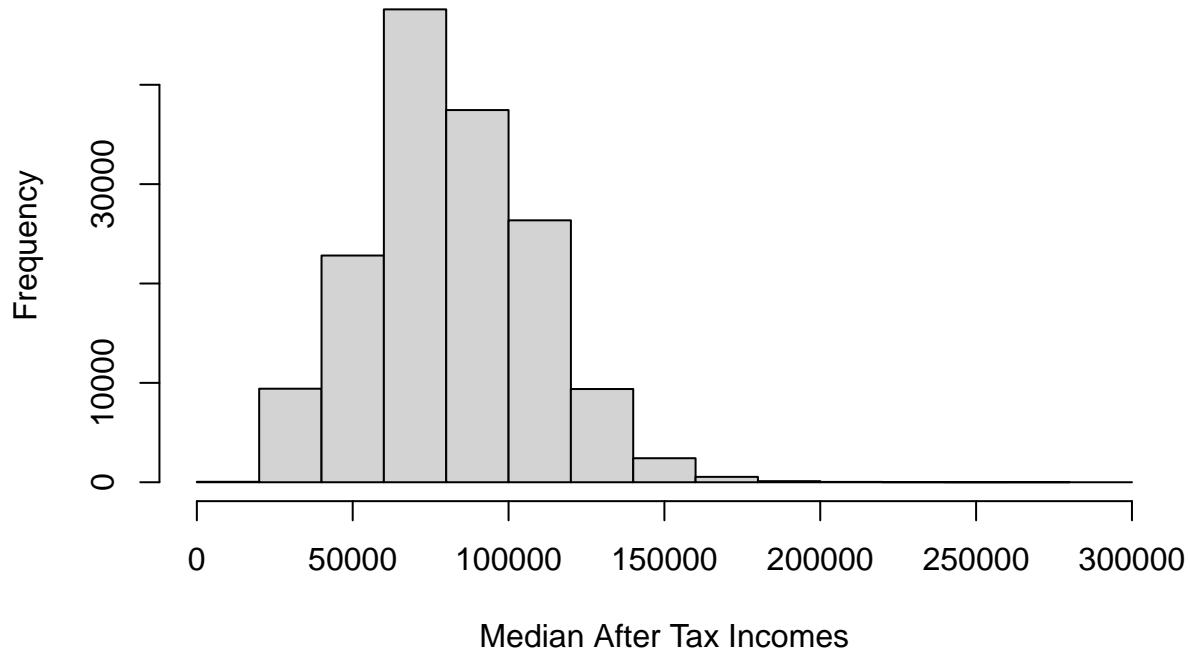
```
box2015 = boxplot(inc2015, ylab = "Median After Tax Incomes", main = "Boxplot of 2015 Median After Tax Income")
```

Boxplot of 2015 Median After Tax Income



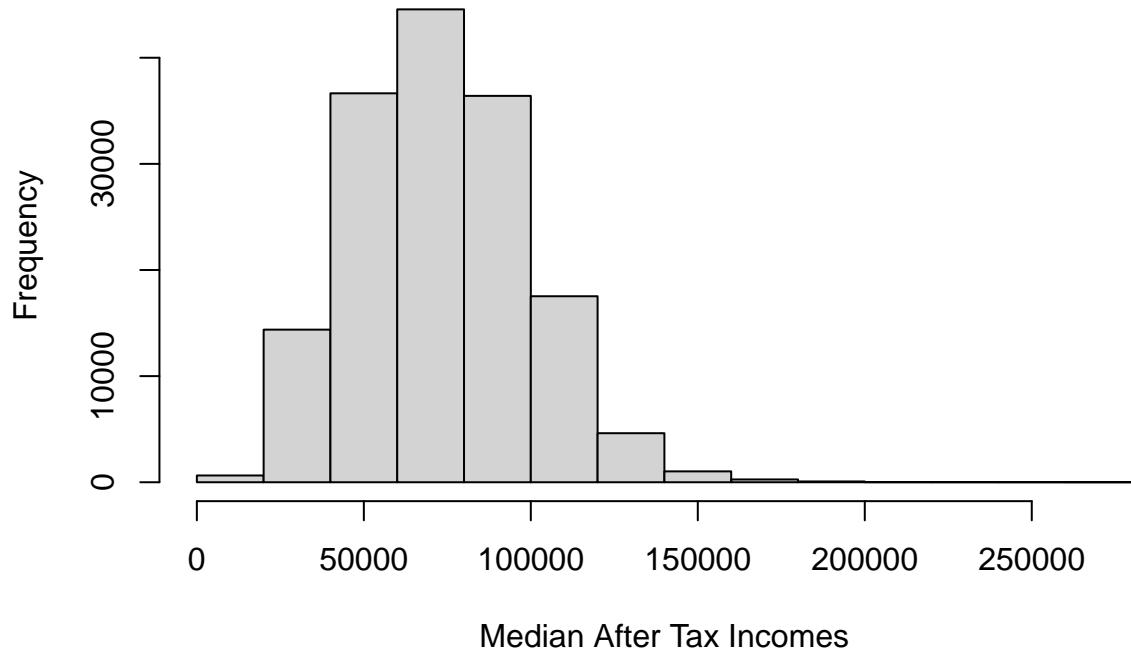
```
# Histograms
hist2020 = hist(inc2020, xlab = "Median After Tax Incomes", main = "Histogram of 2020 median after tax in
```

Histogram of 2020 median after tax incomes



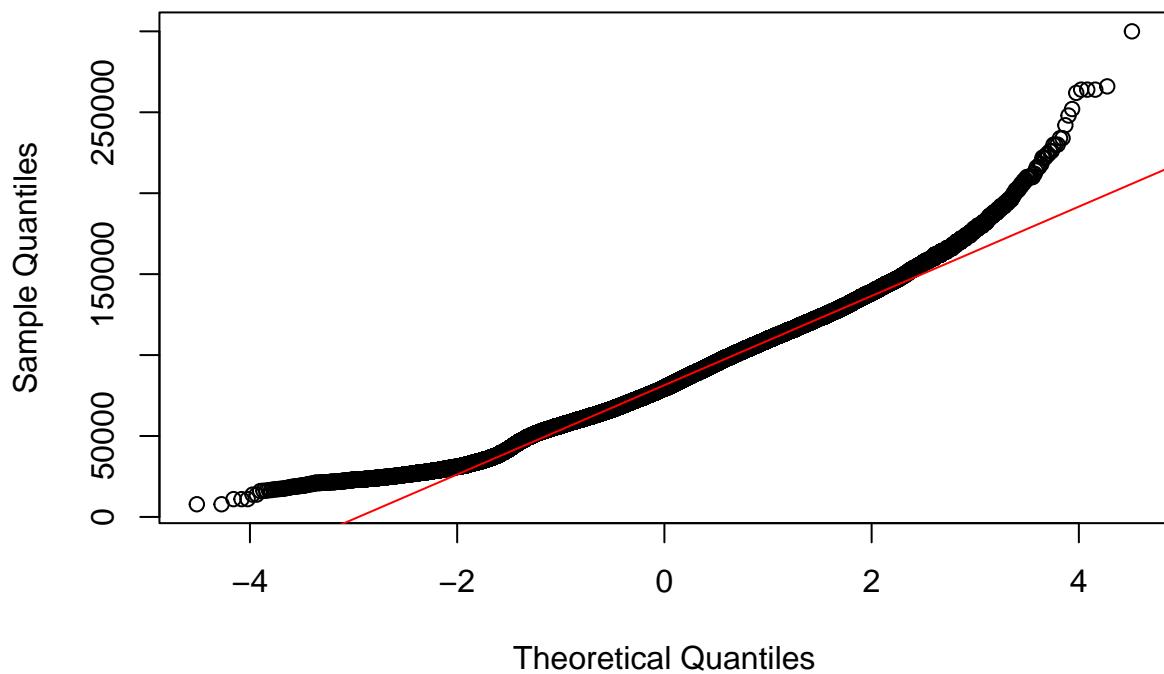
```
hist2015 = hist(inc2015, xlab = "Median After Tax Incomes", main = "Histogram of 2015 median after tax incomes")
```

Histogram of 2015 median after tax incomes



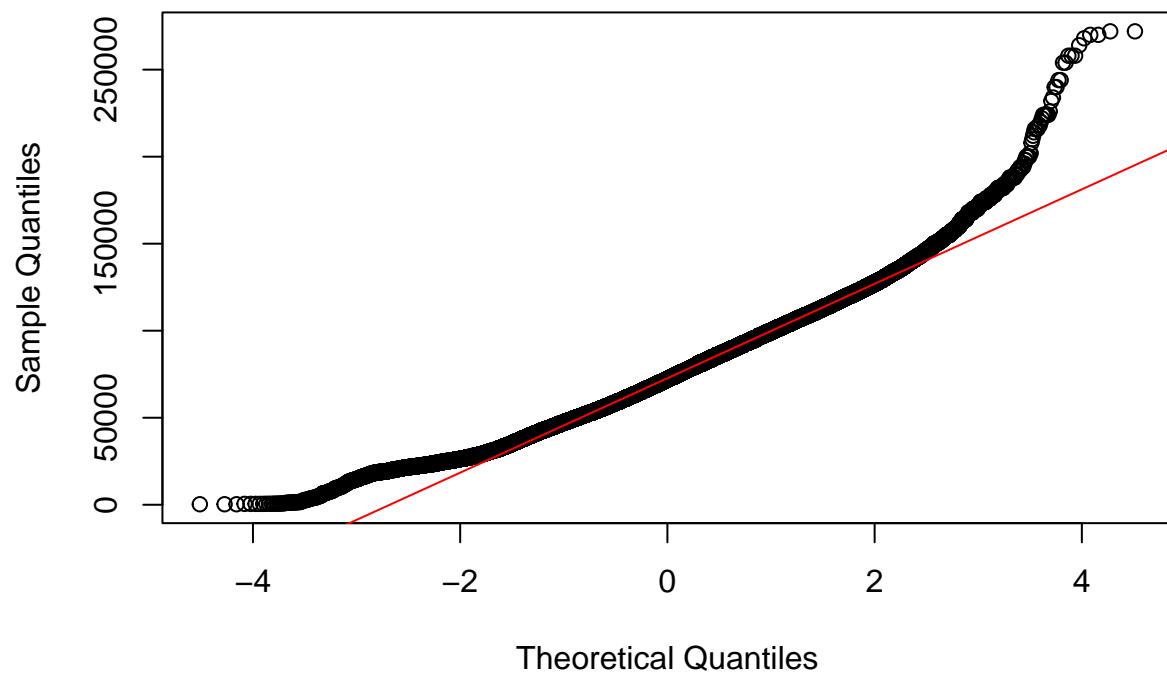
```
# QQ Plots
qq2020 <- qqnorm(inc2020, main = "Normal Q-Q Plot for 2020 Median After Tax Income")
qqline2020 <- qqline(inc2020,col = "red")
```

Normal Q–Q Plot for 2020 Median After Tax Income



```
qq2015 <- qqnorm(inc2015, main = "Normal Q–Q Plot for 2020 Median After Tax Income")
qqline2015 <- qqline(inc2015,col = "red")
```

Normal Q–Q Plot for 2015 Median After Tax Income



Question 2: Is there a significant difference in income between households of 5 or more people and 4 persons (in 2020)?

Load Libraries

```
library(dplyr)
library(ggplot2)
```

Getting data from file

```
income <- read.csv("incomes_cleaned.csv", check.names = FALSE)
```

Data from relevant rows only (rows where household size is ‘4 persons’ and ‘5 or more persons’)

```
income_sub <- income %>%
  filter(GEO != 'Canada') %>%
  filter(
    `Household size (7)` %in% c("4 persons", "5 or more persons"),
    !is.na(`Median household total income (2020)`),
    `Median household total income (2020)` != 0
  )

# Attaching 4 persons household rows
inc4 <- income_sub %>%
  filter(`Household size (7)` == "4 persons") %>%
  pull(`Median household total income (2020)`)

# Attaching 5 or more persons household rows
inc5p <- income_sub %>%
  filter(`Household size (7)` == "5 or more persons") %>%
  pull(`Median household total income (2020)`)
```

Summary statistics, variances

```
var4 <- var(inc4)
var5p <- var(inc5p)

summary(inc4)

##      Min. 1st Qu. Median      Mean 3rd Qu.      Max.
## 26000   108000  124000  122604  139000  274000

summary(inc5p)

##      Min. 1st Qu. Median      Mean 3rd Qu.      Max.
## 34000   111000  128000  128214  145000  428000
```

```
var4  
  
## [1] 807195987  
  
var5p
```

```
## [1] 948878690
```

Two Sample t-test (unequal variances)

T-test since we do not know the population variance

```
test <- t.test(inc4, inc5p, alternative = "two.sided", var.equal = FALSE)  
test
```

```
##  
## Welch Two Sample t-test  
##  
## data: inc4 and inc5p  
## t = -17.111, df = 31036, p-value < 2.2e-16  
## alternative hypothesis: true difference in means is not equal to 0  
## 95 percent confidence interval:  
## -6252.363 -4967.192  
## sample estimates:  
## mean of x mean of y  
## 122603.8 128213.6
```

```
p_val <- test$p.value  
p_val
```

```
## [1] 2.444757e-65
```

Since the p-value is less than 0.05, we reject H₀, and conclude that there is a difference in income between household of 4 persons and households of 5 or more persons

Assumptions

Data is continuous:

Data is independent: Probably true since the two groups are very distinct (4 and 5 or more persons household)

Data is normally distributed: We use boxplot, histogram, and qq plot

```

par(mfrow=c(3, 2))

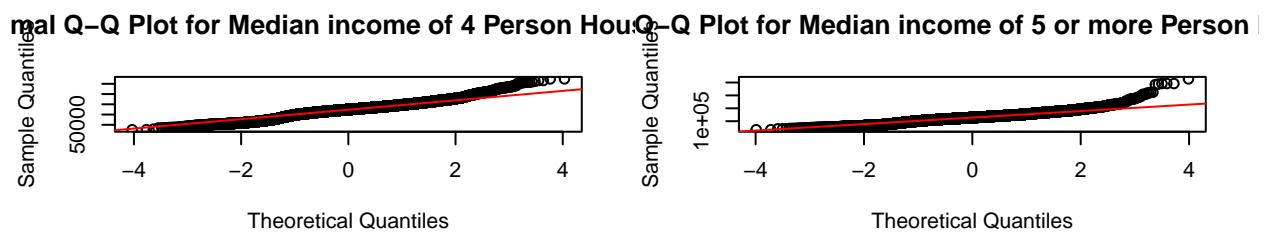
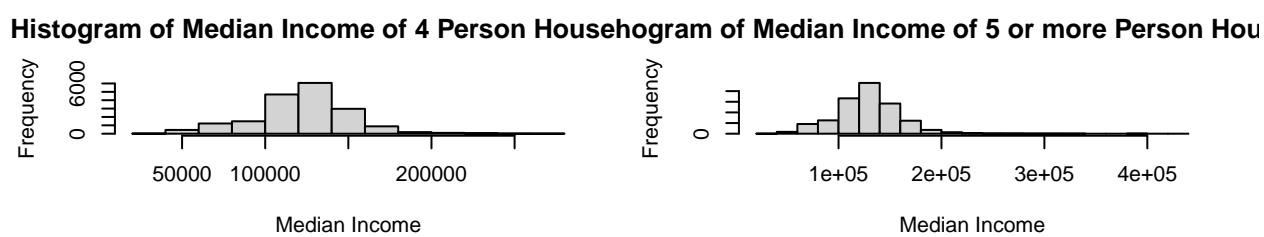
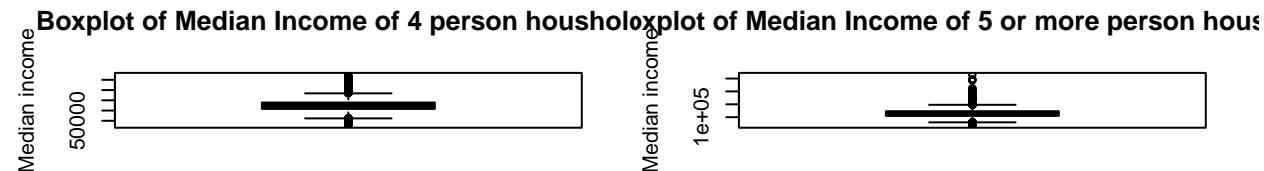
# Boxplots
box4 = boxplot(inc4, ylab = "Median income", main = "Boxplot of Median Income of 4 person households")
box5p = boxplot(inc5p, ylab = "Median income", main = "Boxplot of Median Income of 5 or more person households")

# Histograms
hist4 = hist(inc4, xlab = "Median Income", main = "Histogram of Median Income of 4 Person Households")
hist5p = hist(inc5p, xlab = "Median Income", main = "Histogram of Median Income of 5 or more Person Households")

# QQ plots
qq4 <- qqnorm(inc4, main = "Normal Q-Q Plot for Median income of 4 Person Households")
qqline4 <- qqline(inc4, col = "red")

qq5p <- qqnorm(inc5p, main = "Normal Q-Q Plot for Median income of 5 or more Person Households")
qqline5p <- qqline(inc5p, col = "red")

```



Question 3: Is there a difference in median income between Households with and without children?

Do households with children earn more than households without children?

Load libraries

```
library(dplyr)
library(ggplot2)
```

Getting data

```
income <- read.csv("incomes_cleaned.csv", check.names = FALSE)
```

Filtering data

```
# We check data by the Household type... column
income_sub <- income %>%
  filter(GEO != 'Canada') %>%
  filter(
    # Include the data if the row is 'Without children' or 'With children'
    `Household type including census family structure (11)` %in% c("Without children", "With children"),
    !is.na(`Median household total income (2020)`),
    `Median household total income (2020)` != 0
  )

# Attaching the data
without_children <- income_sub %>%
  filter(`Household type including census family structure (11)` == "Without children") %>%
  pull(`Median household total income (2020)`)

with_children <- income_sub %>%
  filter(`Household type including census family structure (11)` == "With children") %>%
  pull(`Median household total income (2020)`)
```

Summary statistics & variances

```
var_without <- var(without_children)
var_with <- var(with_children)

summary(without_children)
```

```
##      Min. 1st Qu. Median   Mean 3rd Qu.   Max.
##  37600   71000  80000  82240  91000  216000
```

```
summary(with_children)

##      Min. 1st Qu. Median    Mean 3rd Qu.    Max.
##    36400 109000 121000 122755 135000 392000
```

```
var_without
```

```
## [1] 302182460
```

```
var_with
```

```
## [1] 547615912
```

Two Sample t-test (unequal variances)

```
test <- t.test(without_children, with_children, alternative = "two.sided", var.equal = FALSE)
test
```

T-test since we do not know the population variance

```
##
##  Welch Two Sample t-test
##
## data: without_children and with_children
## t = -162.1, df = 26551, p-value < 2.2e-16
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
## -41005.23 -40025.42
## sample estimates:
## mean of x mean of y
## 82239.62 122754.94
```

```
p_val <- test$p.value
p_val
```

```
## [1] 0
```

Since the p-value of our t-test is less than 0.05, we reject H₀, and conclude that there is a difference between the income of households with children and without children

Assumptions

Data is continuous:

Data is independent: This is probably true since people with children wouldn't report that they don't have children when replying to a census, and vice versa

Data is Normally distributed: We will use boxplots, histograms, and qq-plots to check this assumption

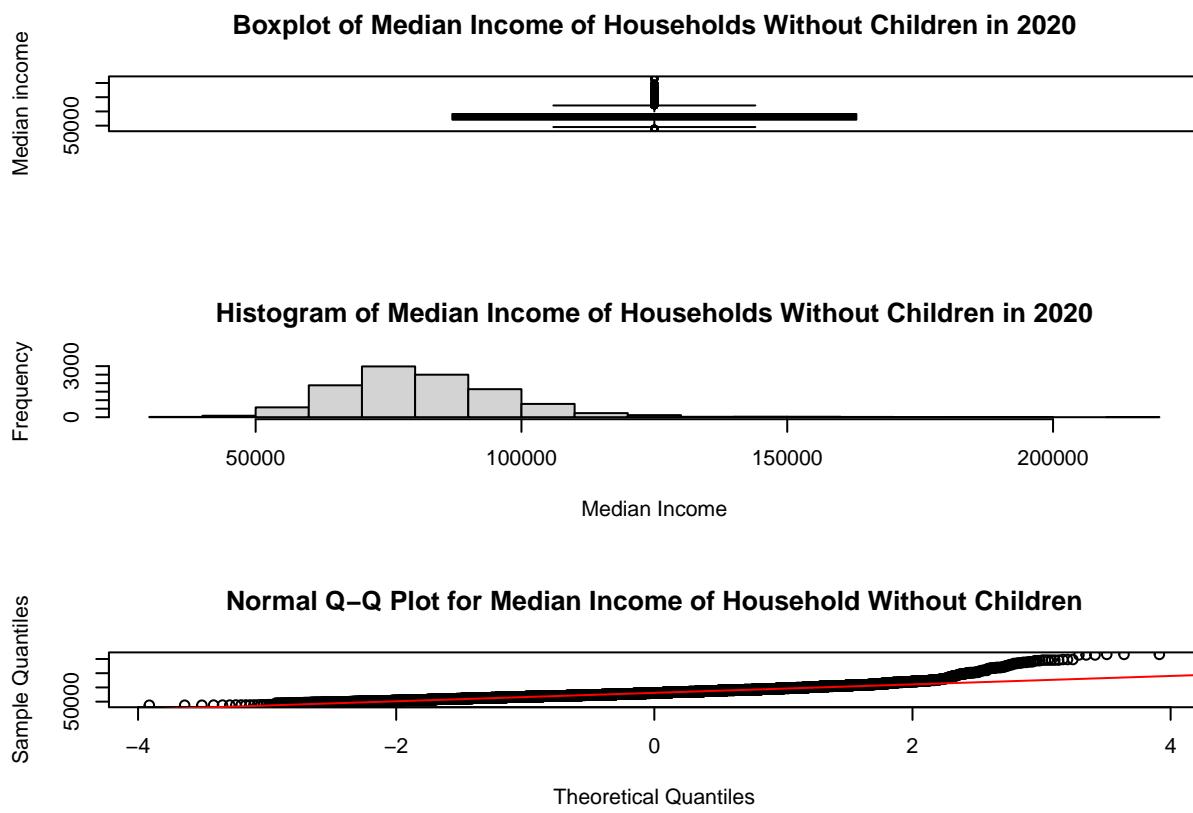
Families without children

```
par(mfrow=c(3,1))

# Boxplot
withoutbox <- boxplot(without_children, ylab = "Median income", main = "Boxplot of Median Income of Households Without Children in 2020")

# Histogram
withoutHist <- hist(without_children, xlab = "Median Income", main = "Histogram of Median Income of Households Without Children in 2020")

# QQ plot
qqwithout <- qqnorm(without_children, main = "Normal Q-Q Plot for Median Income of Household Without Children in 2020")
qqlinewithout <- qqline(without_children, col = "red")
```



Families With Children

```
par(mfrow=c(3,1))
# Boxplot
withbox <- boxplot(with_children, ylab = "Median income", main = "Boxplot of Median Income of Households With Children in 2020")

# Histogram
withHist <- hist(with_children, xlab = "Median Income", main = "Histogram of Median Income of Households With Children in 2020")

# QQ plot
qqwith <- qqnorm(with_children, main = "Normal Q-Q Plot for Median Income of Household Without Children")
qqlinewith <- qqline(with_children, col = "red")
```

